

Santa Clara University
ELEN266
Assignment: Inner Product

Instructor: Hossein Omidian
Copyright: Chris Dick

Winter 2023

In this assignment you will design an FPGA implementation of a circuit to compute inner-products using the Vivado/Vitis High Level Synthesis (HLS) design flow.

This assignment counts towards your final grade and is worth 20 marks, instructions for submitting your work are provided at the end of this document.

Some design files will be supplied as a starting point for the project, this will be discussed in class.

Note that some of the GUI screenshots in this document might not look exactly like the corresponding GUIs in the version of the tools you are using, but that's ok, if there are any differences, they will be relatively minor¹

1 Inner-Product

The inner product $\langle a; b \rangle$, also referred to as dot product, of two length N column vectors a and b is defined as

$$\langle a; b \rangle = a^T b = \sum_{i=0}^N a_i b_i \quad (1)$$

Evaluating Equation 1, requires N multiplications and $N - 1$ additions, so it is a calculation that is of order N , or in big O notation, $O(N)$. Note that N memory reads are required to read all of the values in the a and similarly N memory reads are required to access all of the data in the b array. One write operation is required to return the final result to memory.

The inner product is a very common linear algebra operation that forms the foundation of many machine learning and signal processing functions such as filtering and matrix multiplication.

In this assignment you will learn how to implement inner product calculations on an FPGA using C as the programming language. You will learn how to

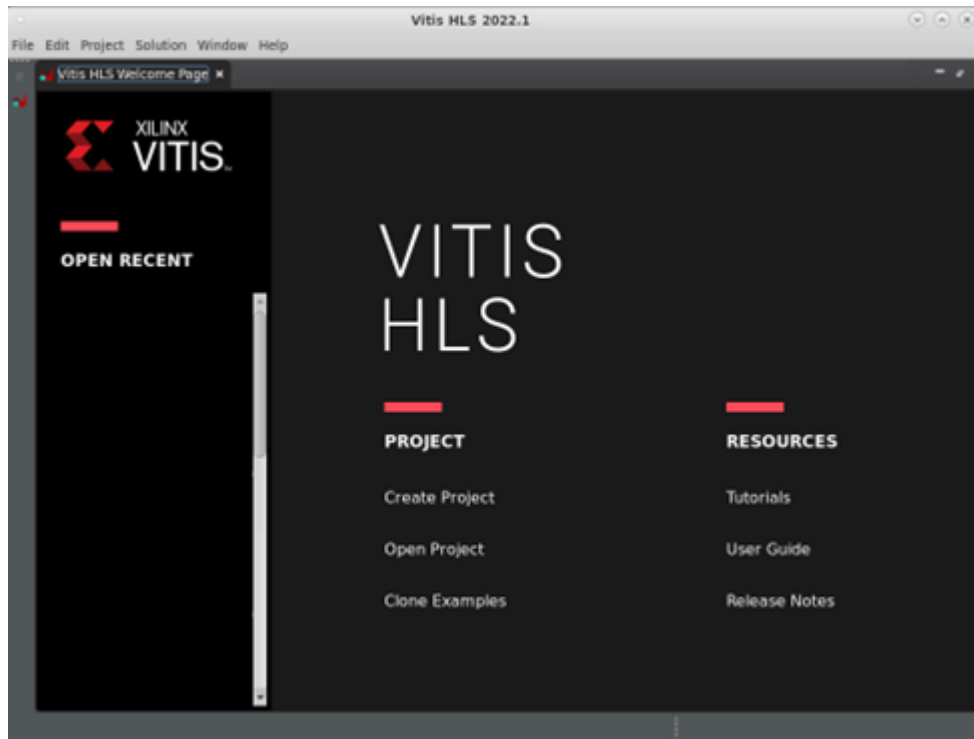


Figure 1: Vitis HLS welcome screen

use the AMD Vivado/Vitis HLS tool chain and how to generate different versions of $\langle a; b \rangle$ that provide design tradeoffs between hardware cost (area) and performance as measured by the number of inner product calculations computed per second.

1.1 Baseline Implementation: Solution 1

This section describes the process for generating an HLS implementation of the inner product calculation.

The following sections of this document contain screenshots of various Vitis HLS windows. They may not exactly match the screens that you will see when running the tools as there are some differences depending on the version of the tools that you are running.

Start Vitis HLS. The Welcome screen of Figure 1 will open. Create a new project by selecting Create New Project located on the left of the window. In the next few steps the C/C++ source code for the design and associated testbench and other files required by the testbench are added to the project.

The top-level design function, the function that is to be implemented in the FPGA is called *inner_product*. The corresponding source file is *inner_product.cpp*.

In the text box at the top of Figure 2 enter the top-level function name, that is, *inner_product*. Click Add Files, navigate to the location of the file *inner_product.cpp*, select it, and add it to the project. If the design employs multiple C/C++ source files, then all of the files would be added to the project. This project only has one source code file. Click through to the next configuration screen by selecting Next located at the bottom of the window. Click Add Files on the page shown in Figure 3 to add the testbench *tb_inner_product.cpp* to the project. The testbench has a dependency on the header file *inner_product.h*. The header file is automatically added to the project as long as it resides in the main project directory, that is the directory containing *inner_product.cpp* and *tb_inner_product.cpp* in this case. Click through to the Solution Configuration page shown in Figure 4. Leave the default name, Solution1 in the Solution Name textbox. The Clock Period textbox captures the target clock period for the design, the target operating clock frequency for the FPGA. There is no guarantee that the design will, out-of-the-box, achieve this target. It might be necessary to apply various implementation directives and re-factor the C/C++ source code in order to close timing and accomplish other design objectives, particularly for high clock frequency designs. The next step is to select the FPGA device by clicking through the button on the right side of the Part Selection section of the page. From the menu shown in Figure 4 select the *Kintex xc7k160tbg676-2* device.

Click Finish and the Vitis HLS Integrated Development Environment (IDE) screen shown in Figure 6 is displayed. Some annotations have been overlaid on this page to identify the functions for several of the key icons in the menu ribbon.

Go to *Solution* in the menu ribbon and select *Solution Setting*. In the *General* tab, select *config.compile* and make sure *pipeline_loops* Value is set to zero as it's shown in figure 7.

The IDE provides an environment for editing code, executing C/C++ code, RTL simulation and running the entire FPGA tool chain starting with high-level synthesis of the C/C++ source code through to place-and-route in the backend physical implementation tools. The project files can be accessed from the Explorer panel on the left of the page. Open *inner_product.cpp* from the Explorer panel by clicking on source and selecting *inner_product.cpp* from the file list - in this case there is only a single file in the list.

Listing 1 show the testbench for the project. Two length N vectors, a and b are declared (line 13), these will be the inner product input vectors. Variable s will hold the result of the inner product calculation produced from the golden reference inner product function and $s1$ will store the result of the inner product computed by the design function described in *inner_product.cpp*. The input arrays are initialized with some random integers with the for loop starting on line 15. Line 21 is the call to the inner product reference function *inner_product_reference()* and line 21 calls the design function *inner_product()*. Line 26 validates the design function by comparing the result of the inner product, $s1$, generated by the design function with the golden reference value s . The comparison sets the simulation status flag, *sim_status*, to 1 if a failure is




Figure 2: Adding the C/C++ source code to the project




Figure 3: Adding the project testbench.

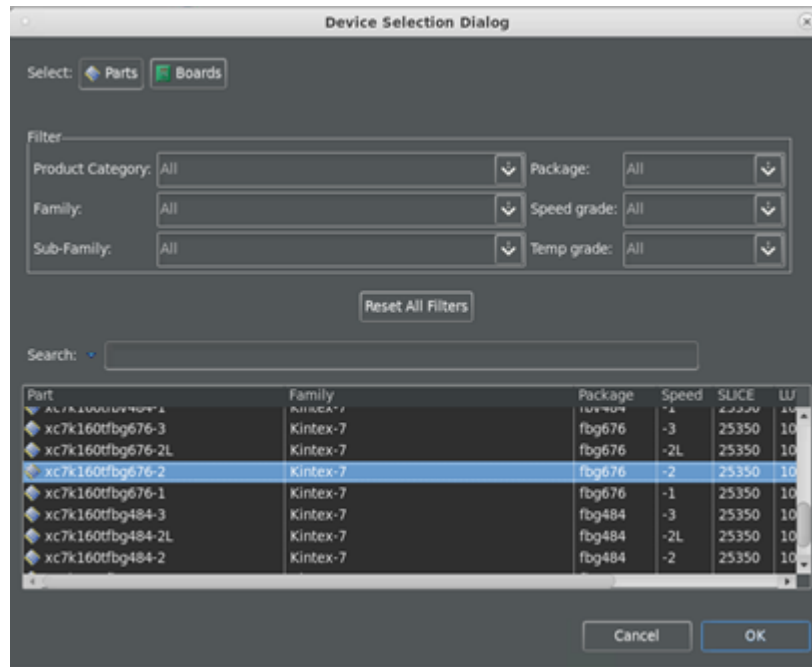


Figure 4: Selecting the target FPGA

detected. Line 36 returns the result of the simulation to the caller. It is best practice to always return a value to the caller of `main()`. A return value of 0, or false, indicates a successful execution, while a return value of 1, or true designates a failure to the caller.

Listing 1: Testbench for the inner product design.

```

1  #include <math.h>
2  #include <iostream>
3  #include <fstream>
4  #include "inner_product.h"
5
6
7  int main ()
8  {
9      int sim_status = 0;
10
11     int i;
12
13     data_t a[N], b[N], s, s1;
14
15     for (i=0; i<N; i++)
16     {
17         a[i] = rand() % MAXNUM;
18         b[i] = rand() % MAXNUM;
19     }

```

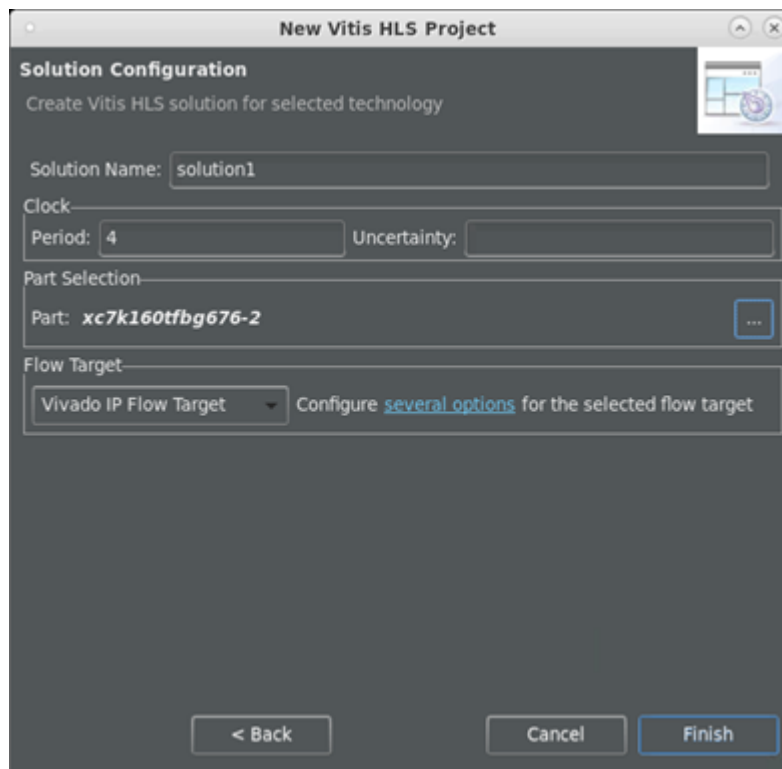


Figure 5: The final page in the project configuration setup

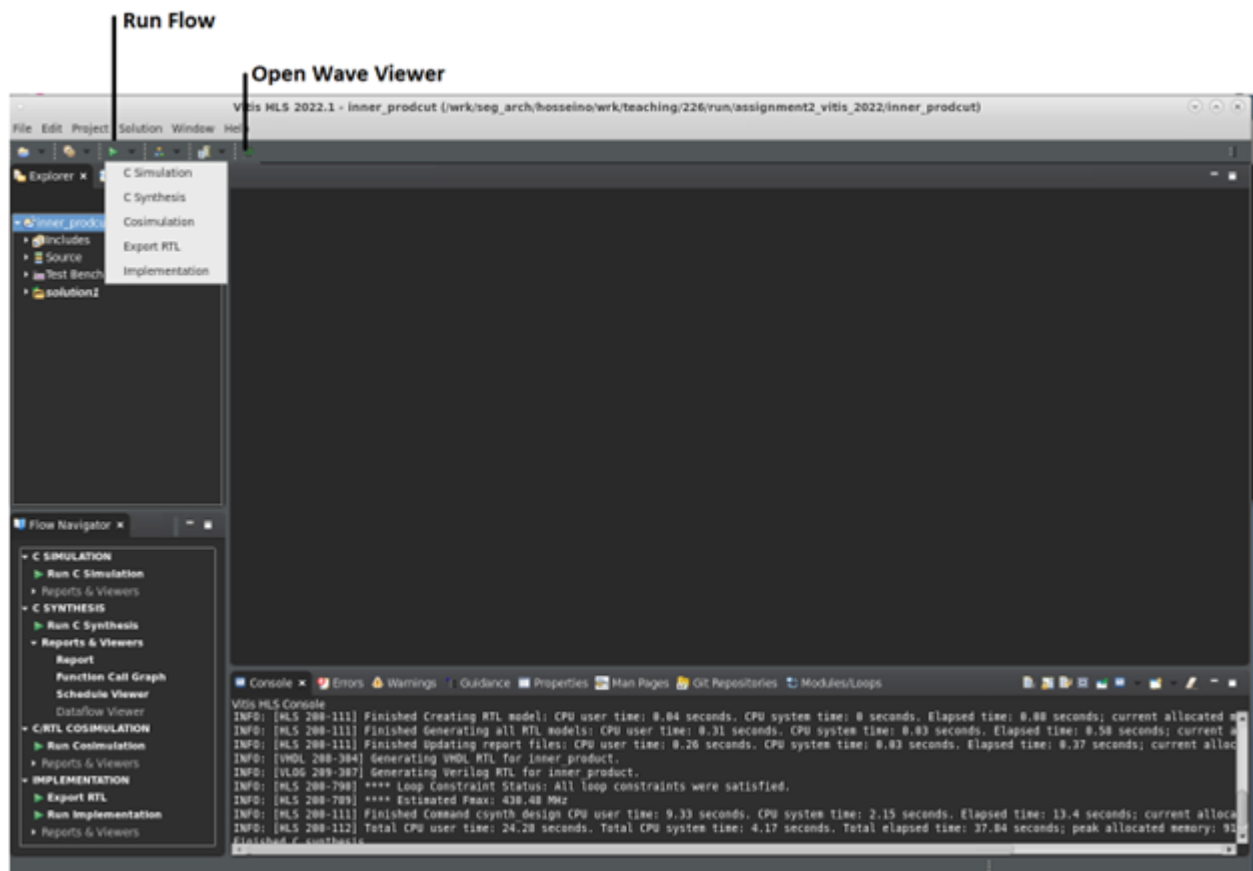


Figure 6: Vitis HLS IDE

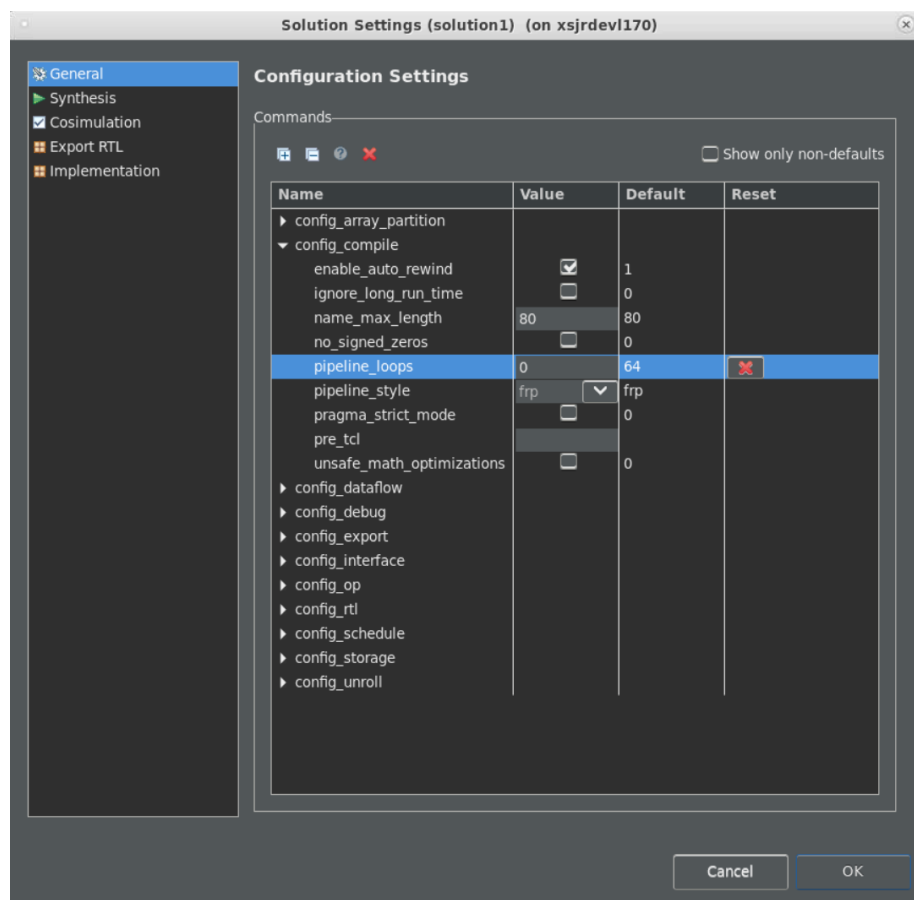


Figure 7: Solution Settings - Pipeline setting

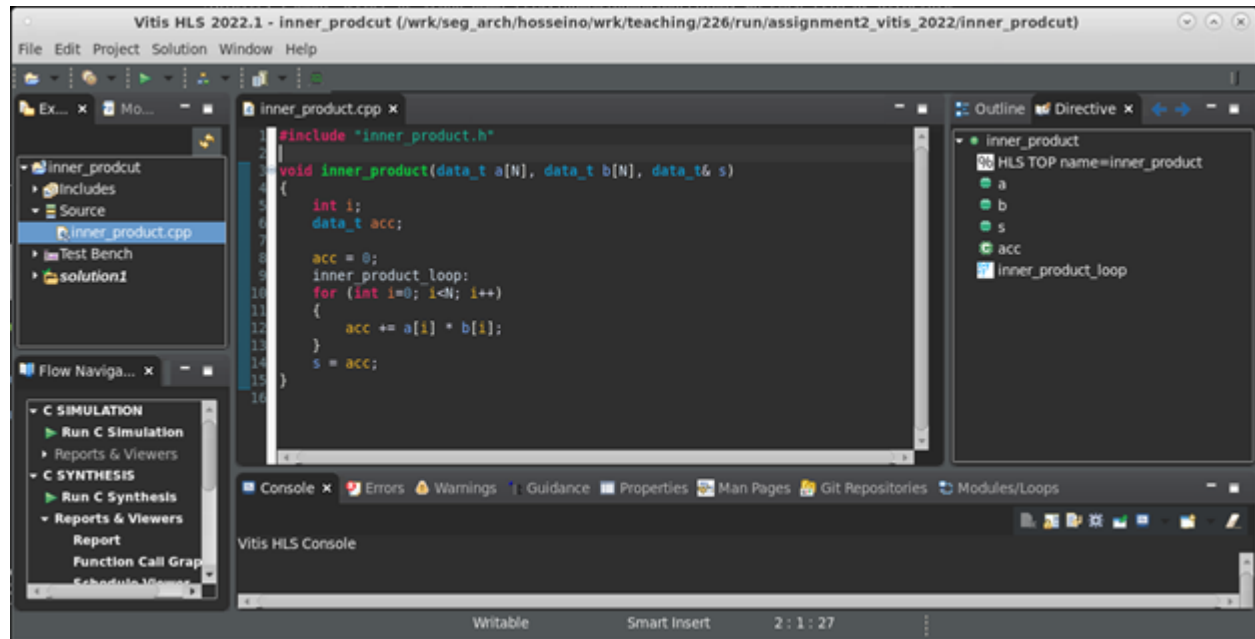


Figure 8: Opening the C++ source file *inner_product.cpp*.

```

20
21     inner_product_reference(a, b, s);
22     inner_product(a, b, s1);
23
24     std::cout << s << '\t' << s1 << std::endl;
25
26     if (s != s1)
27         sim_status = 1;
28
29         // clean-up
30
31     if (sim_status==0)
32         std::cout << "simulation passed" << std::endl;
33     else
34         std::cout << "simulation failed" << std::endl;
35
36     return(sim_status);
37 }
38
39 void inner_product_reference(data_t a[N], data_t b[N], data_t &s)
40 {
41     s = 0;
42     for (int i=0; i<N; i++)
43         s += a[i] * b[i];
44 }
45

```

Listing 2: Header file for inner product design project.

```

1 #ifndef _INNER_PRODUCT_H_
2 #define _INNER_PRODUCT_H_
3 #include "ap_fixed.h"
4 #define N 8
5 #define MAXNUM 100
6
7 typedef int data_t;
8
9 // function prototypes
10 void inner_product(data_t a[N], data_t b[N], data_t &s);
11 void inner_product_reference(data_t a[N], data_t b[N], data_t &s);
12
13 #endif

```

Listing 3: C source code expressing the required inner product calculation.

```

1 #include "inner_product.h"
2
3 void inner_product(data_t a[N], data_t b[N], data_t& s)
4 {
5     int i;
6     data_t acc;
7
8     acc = 0;
9     inner_product_loop:
10    for (int i=0; i<N; i++)
11    {
12        acc += a[i] * b[i];
13    }
14    s = acc;
15 }

```

Listing 4: HLS Console window pane reporting the steps invoked when running a C simulation.

```

1 INFO: [SIM 2] ***** CSIM start *****
2 INFO: [SIM 4] CSIM will launch GCC as the compiler.
3     Compiling ../../../../inner_product_source_files_v1/
4         tb_inner_product.cpp in debug mode
5     Compiling ../../../../inner_product_source_files_v1/
6         inner_product.cpp in debug mode
7     Generating csim.exe
8 29111 29111
9 simulation passed
10 INFO: [SIM 1] CSim done with 0 errors.
11 INFO: [SIM 3] ***** CSIM finish *****

```

The header file *inner_product.h*, Listing 2 contains the definitions for the data types used in the design and also the length of the vectors *a* and *b*.

Listing 3 is the C code for the design function *inner_product()*, again this is the compute kernel that is to ultimately be realized in the FPGA. The testbench, Listing 1, is purely for design verification and is not compiled to hardware.

Simulate the design by clicking on the Run C Simulation button in the IDE Run Flow drop-down, refer to Figure 6. The dialog window shown in Figure 9

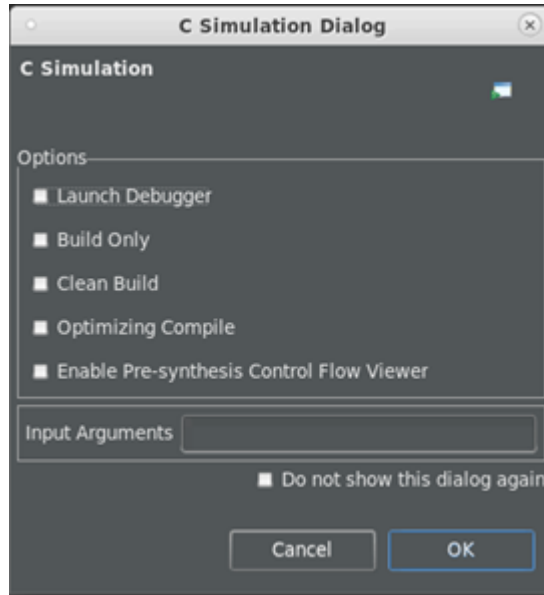


Figure 9: This dialog window is opened when first running C simulation from within the HLS IDE.

will open. Referring to Figure 10, check the Clean Build and Do not show this dialog box again checkboxes and then click OK to start the simulation. Figure 11 shows the simulation dialog in the Console window. For clarity, the contents of the Console pane are shown in Listing 4. Read through the summary and note the points in the flow where the various project files are opened, the target clock period and FPGA device for the design reported, and in the final part of the dialog confirmation of the final result (pass or fail) is reported.

After a successful C simulation C synthesis can be invoked (refer to Figure 6 for the icon in the IDE Run Flow drop-down to select) to generate the RTL (Register Level Transfer Logic) for the design. Figures 12 to ?? show the synthesis report for the design.

Figure 12 shows the timing and latency summary for the design. The target clock frequency was defined as 4 ns, high-level synthesis is reporting that the estimated clock frequency of the final design is 2.323 ns, which meets the timing goal with a margin of $4 - 2.323 = 1.08$ ns. Note that this is an estimate, the final operating frequency of the design can only be conclusively known once the design is run through place-and-route and timing analysis performed in this final design. Figure 12 also summarizes the design resource utilization. Figure 13 shows the block interface definition. This is the design hardware interface for this component.

Selecting the Open Viewer tab from the menu ribbon reveals the design perspective shown in Figure 14. As we can see Iteration Interval (II) for the

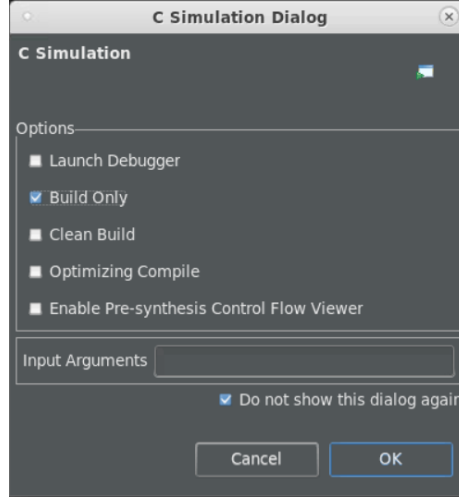


Figure 10: C simulation dialog window with several simulation options selected.

design is $II = 5$ clock cycles and that the trip count is $N_{trip} = 8$. The iteration latency is the latency, specified in clock cycles, for completing one iteration of the main compute loop, *inner_product_loop* in Listing 3. The trip count is the number of times the loop is executed, which is equal to the length of the vectors **a** and **b**. Recall that this example was compiled with $N = 8$. Since this loop is not pipelined, iteration $i + 1$ does not begin until iteration i has completed. Therefore, the latency of the compute loop is

$$L_{compute} = II \times N_{trip} = 5 \times 8 = 40 \quad (2)$$

Referring again to the Performance Profile, observe that the function level latency is $N_{func} = 41$ clock cycles, which means that it takes 41 clock cycles to deliver the result of the calculation to memory. 40 clock cycles correspond to compute and then an additional cycle to write the result to memory. The Initiation Interval for the function is 42 clock cycles. This defines the minimum time interval between successive calls to the function, or the time between consuming successive function input operands. In this case one $N = 8$ length inner product can be performed at the rate of one vector operation per 42 clock cycles. The majority of this time is composed of the compute latency $L_{compute} = 40$.

For a design with a single Multiply-Accumulator (MAC), the theoretical lower bound for computing the length N inner product is N clock cycles. Or expressed another way, one vector operation can be computed every N clock cycles. This of course is supported with a new MAC operation every clock cycle, corresponding to $II=1$. In the next section of the assignment, we will see how to provide architectural design directives to realize a time-optimal design with $II=1$.

C synthesis generates an RTL representation of the design. The next step

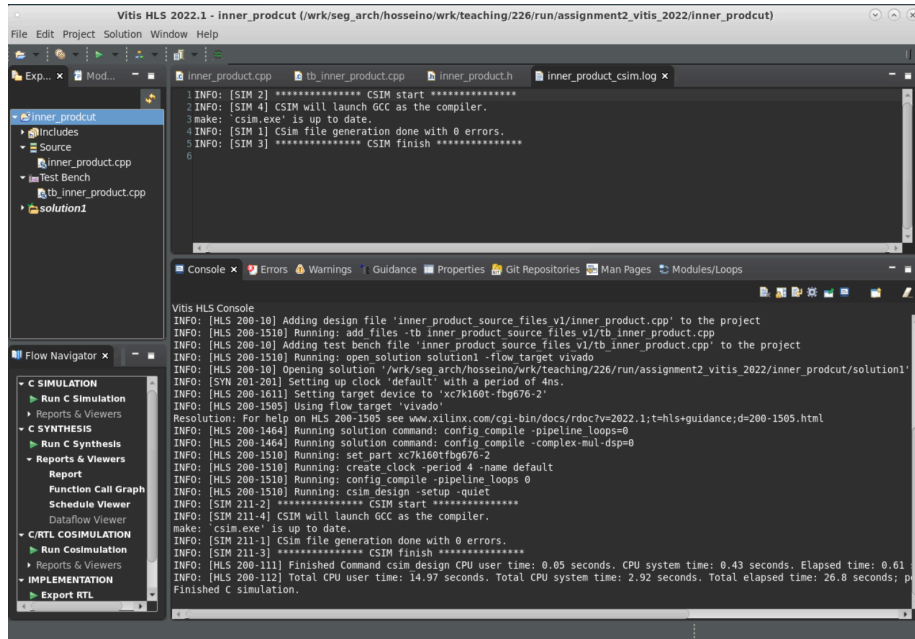


Figure 11: The HLS Console panel (the bottom pane in the Figure) provides status information on the various steps involved in the simulation and reports the simulation completion status, in this case the pane is indicating that the simulation Csim was completed with 0 errors

in the design flow is to validate the functional correctness of the RTL with the original generative C code. Theoretically, since the RTL is produced from the C code it should be correct by construction, and most of the time that is the case. However, it is possible, through the specification of certain pragmas or implementation directives, to produce a design that is not functionally correct. This can occur if the programmer makes incorrect assumptions about, for example, memory access behavior of the design, and then, these design directives, provide incorrect guidance to the C synthesis flow so causing synthesis to generate incorrect RTL. There is also always the possibility, and of course this is true of any software, that the C synthesizer and related tools has bugs, and that this could be the cause of an error in the C synthesis process. So in general, it is considered good practice to cross verify the RTL with the C source code. This step is referred to as Run C/RTL Cosimulation in the Vitis HLS flow, refer to Figure 6.

Starting the cosimulation process will open the window shown in Figure 15. Select the Vivado XSIM from the pull-down sub-menu in the Verilog/VHDL Simulator section. Select the **all** option from the Dump Trace pull-down. This will make all of the signals in the design available to the waveform viewer that we will look at next. All other options can be left at their default values. Now click

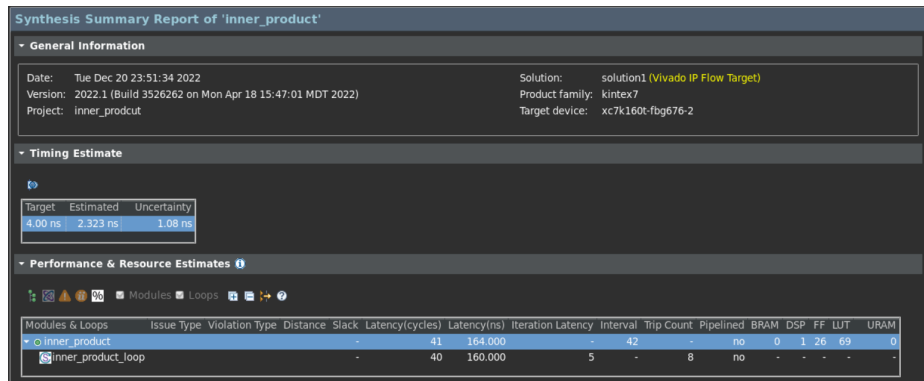


Figure 12: C synthesis report - top part of IDE Synthesis Report panel.




Figure 13: Hardware Interfaces report

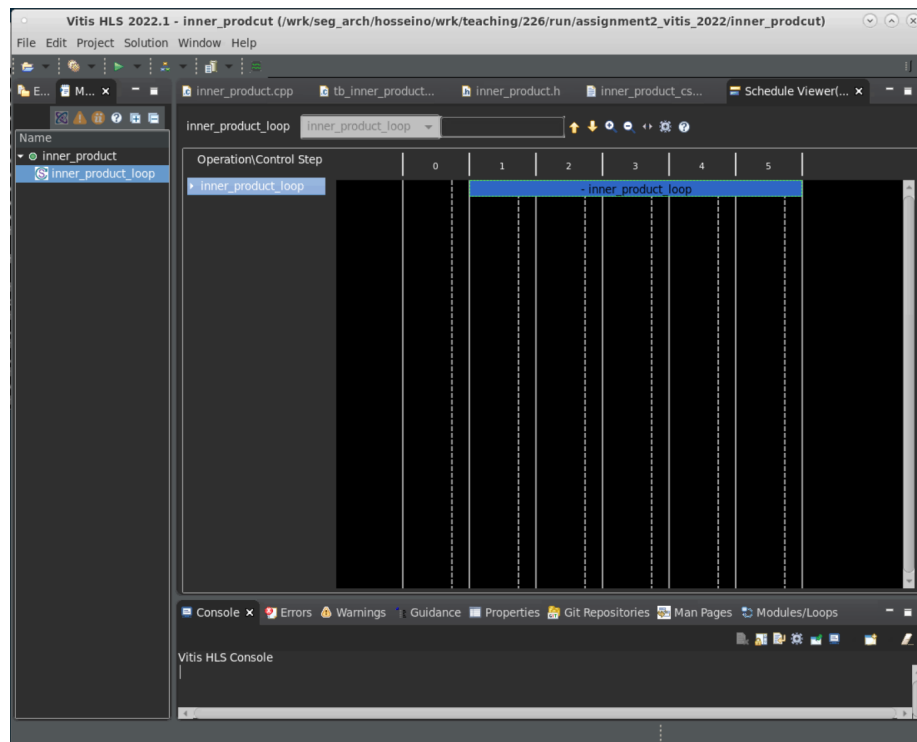


Figure 14: Caption

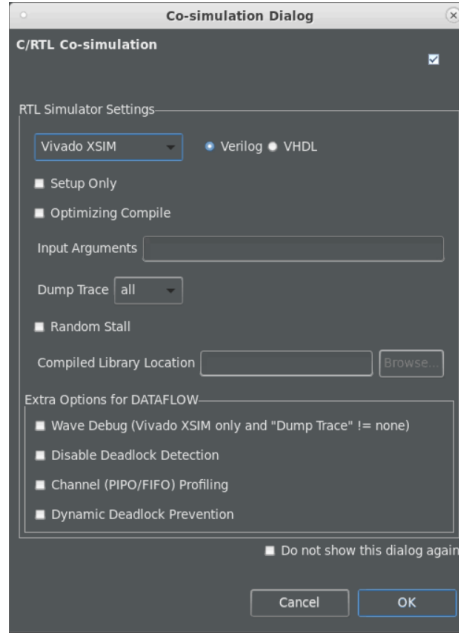


Figure 15: Running and configuring C/RTL cosimulation.

OK to start cosimulation. Figure 16 shows the console status for a successful cosimulation.

After RTL cosimulation has been completed the RTL simulation database can be viewed in the Vitis Waveform Viewer that is invoked using the Open Wave Viewer in the HLS IDE menu ribbon. The resulting waveform is shown in Figure 17.

In this simulation the input vectors are initialized as $a = [83, 77, 93, 86, 49, 62, 90, 63]$ and $b = [86, 15, 35, 92, 21, 27, 59, 26]$. The running sum of the point-by-point products that lead to the final inner product is as follows

$$7138, 8293, 11548, 19460, 20489, 22162, 27473, 29111 \quad (3)$$

These sequences can be observed in Figure 17. Also shown is the start signal being asserted and the `ap_ready` and `ap_done` signals. The transition of start to true (logic 1) initiates the calculation. We see the memory addresses for the memories storing `a` and `b` being generated and associated data values being read from the memories. When `ap_done` is asserted, as shown in the lower part of the figure, the inner product calculation is available. This happens 41 clock cycles after start is asserted. On the 42nd clock cycle the next inner product calculation begin.

The next step in the flow is to place-and-route (PAR) the design, this step is initiated via Implementation on the menu ribbon, refer to Figure 6. After selecting Implementation the window in Figure 18 is presented. Select the Vi-

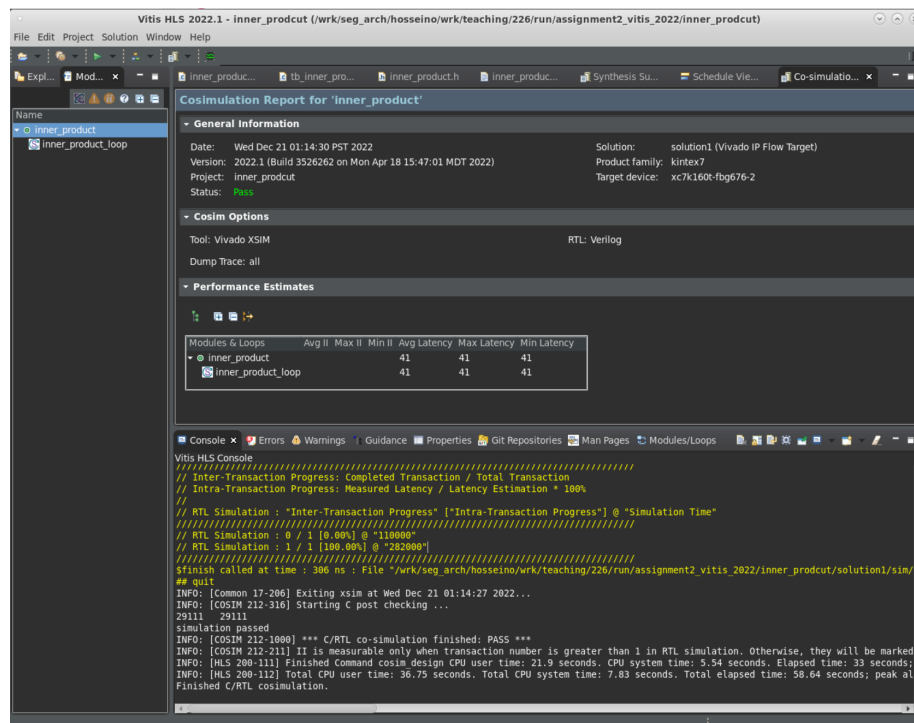


Figure 16: C/RTL cosimulation console dialog after succesful completion of the process.

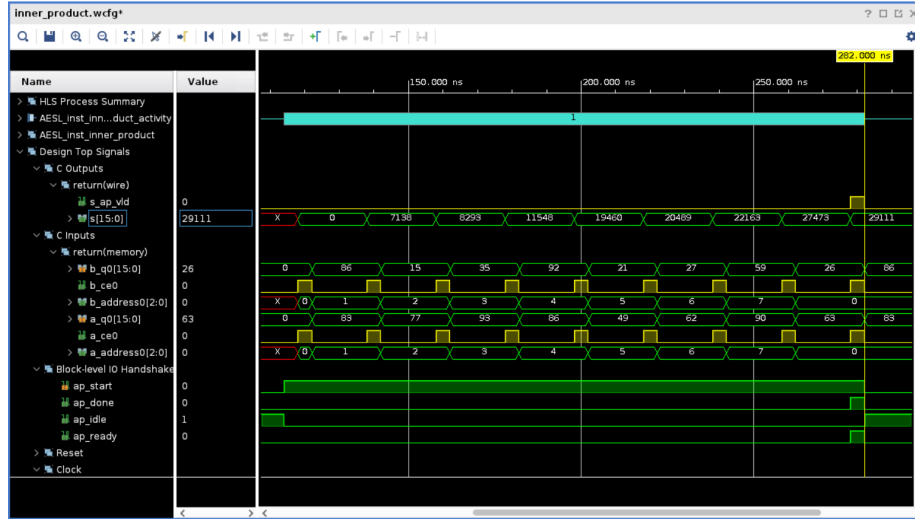


Figure 17: Timing diagram for the $N = 8$ inner product design. The view presented is generated using the Vivado simulator, automatically invoked by selecting Open Wave Viewer from the HLS IDE.

vado synthesis, synthesis, place route check box and click OK. This will invoke the FPGA physical implementation tool flow, and execute the steps of logic synthesis and place-and-route.

Figure 19 shows the Report and console dialog after completion of the physical implementation flow.

1.2 Loop Pipelining: Solution 2

In this section of the assignment a 2nd design implementation, or *Solution*, will be produced using the same C source code of Section 1.1, but with guidance directives supplied to high-level synthesis to produce a throughput optimized implementation. The goal is to generate a design that employs a single multiplier, just like in Solution 1, but that has a much greater efficiency by engaging the multiplier for useful work on every clock cycle. That is, the loop is to have $\Pi=1$. This is accomplished by *loop pipelining* in which iteration $i+1$ starts on the clock cycle following the start of iteration i .

Create a new solution by selecting Project/New Solution from the menu. Leave the default values in the Solution Configuration window, click Finish.

Bring the *Directives* sub-window in to scope as shown in Figure 20. To pipeline the *inner_product_loop* right click on the *inner_product_loop* label in the Directives window to open the directives menu shown in Figure 21. Selection Pipeline from the Directive pull-down . Click ok and the Directives window will look like Figure 22. Note the HLS PIPELINE rewind directive that is attached to the loop *inner_product_loop*. Run C synthesis and review the performance

Run Implementation

Run Vivado RTL Synthesis/Place & Route

☐ RTL Synthesis
☒ RTL Synthesis, Place & Route

RTL

Clock Period

Generate DCP ☐

IP Location

Vivado IP

IP OOC XDC File

IP XDC File

RTL Synthesis Options

RTL Synthesis Strategy

Synth Design Arguments

Place & Route Options

Run Physical Optimizations

Implementation Strategy

Report Options

Report Level

Max Timing Paths

☐ Do not show this dialog again

Figure 18: Running implementation (Place-and-Route)

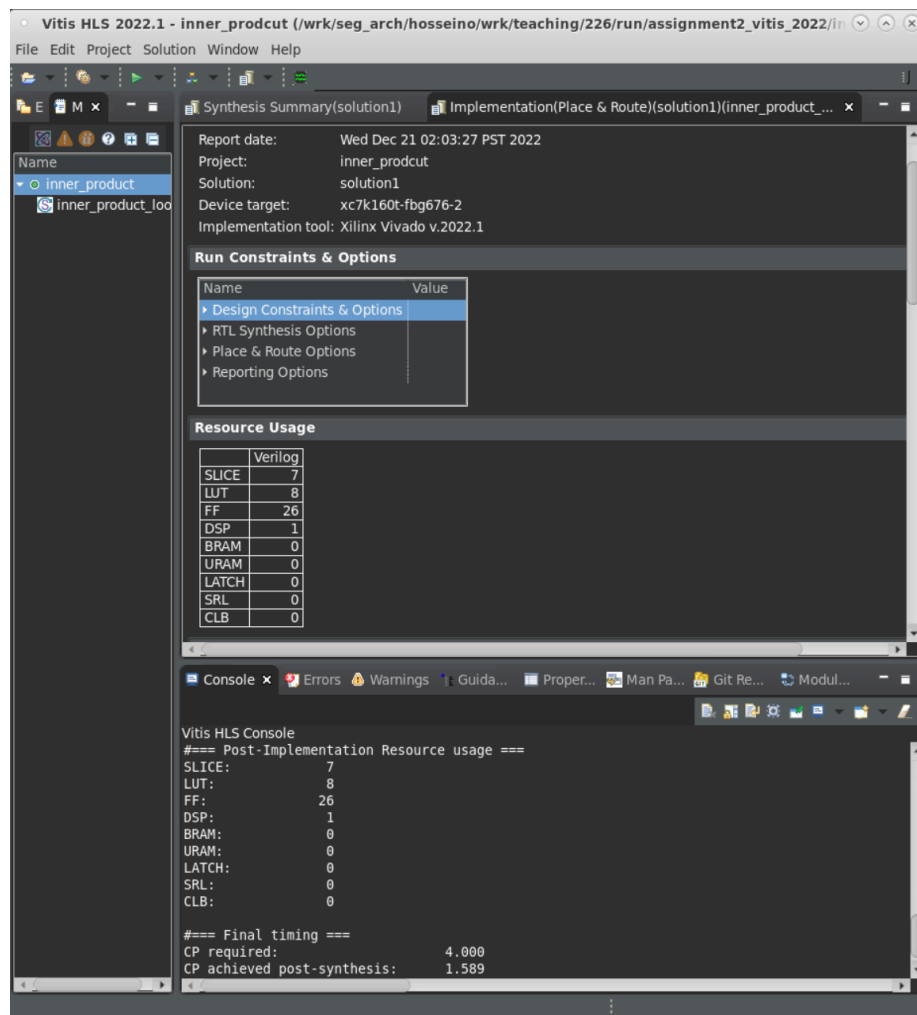


Figure 19: The Report and Console dialog after completing design Implementation

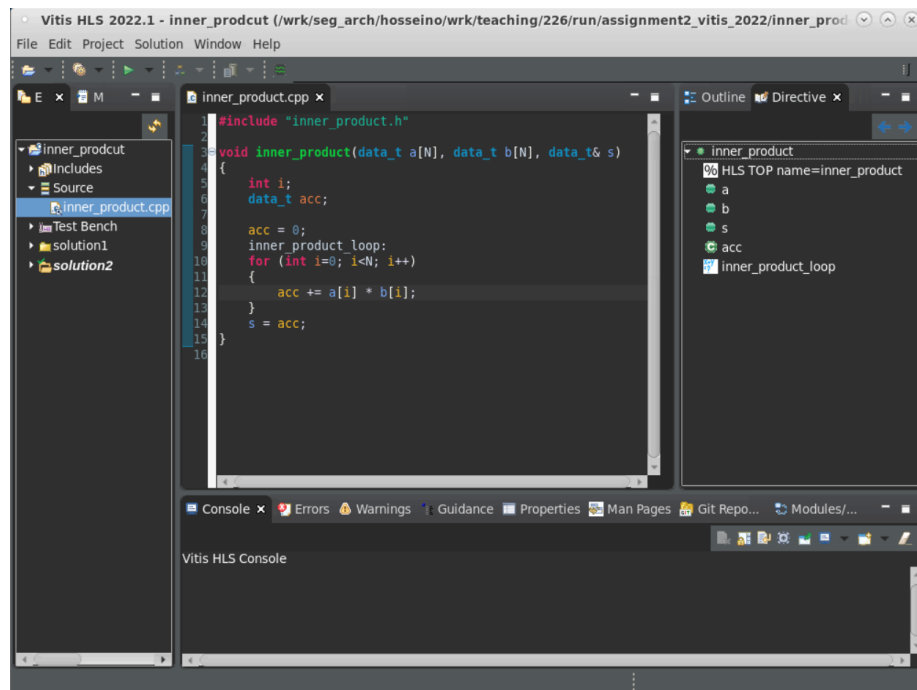


Figure 20: With the *Directives* panel in scope, implementation directives can be added to the design to generate different performance-cost points in design space.

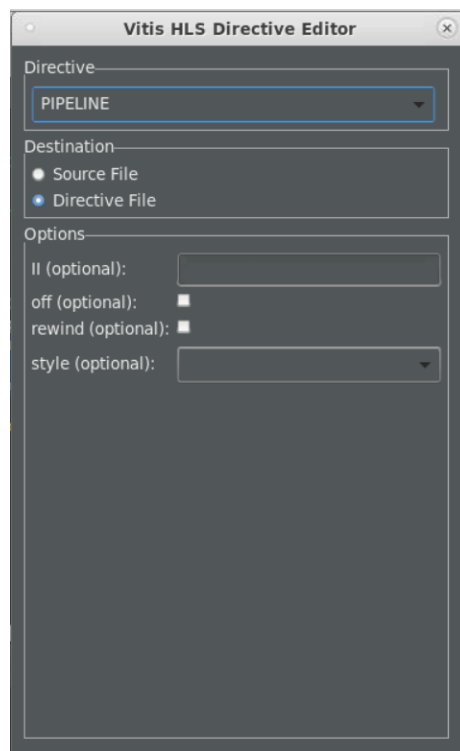


Figure 21: The *Directives* menu.

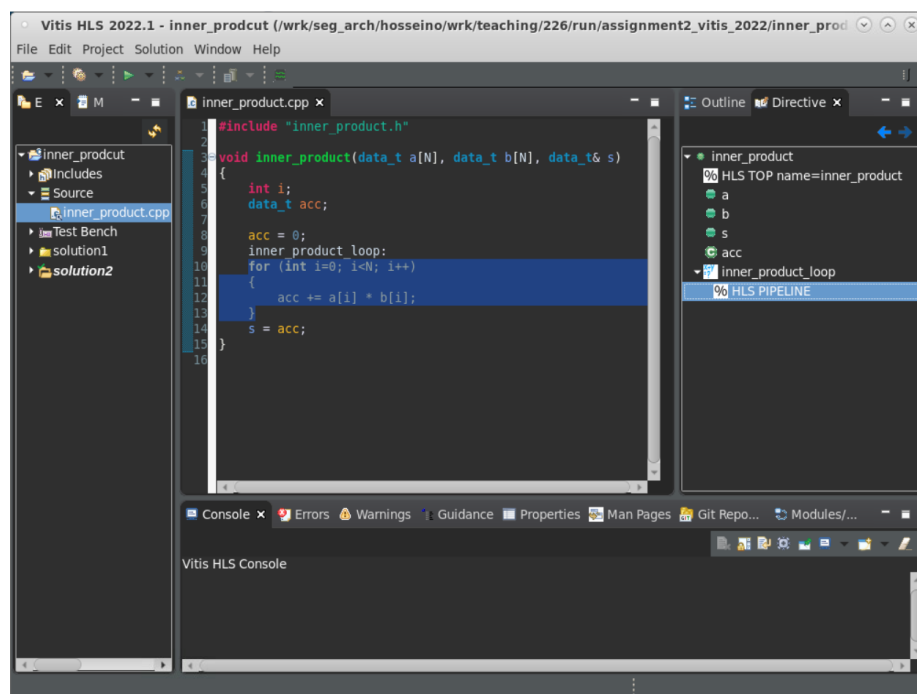


Figure 22: The *Directives* window showing that the *inner_product_loop* and associated pipeline directive.

profile in the Analysis context. Note that the loop has an $II=1$ and the function has $II=14$. With a loop II of 1 a new multiplication is initiated on every clock cycle. The function initiation interval of 8 indicates that the function can be invoked every 8 clock cycles, which is a cadence equal to the length of the input vectors in this example.

Contrast the loop II of 1 in this Solution compared to the II of 4 in the first design. In the first design the multiplier was engaged 25% of the time. In Solution 2 the multiplier is operating at 100% efficiency. For an FPGA clock frequency of f_{FPGA} MHz, the 2nd design is realizing f_{FPGA} Multiply-Accumulates per second (MACs) compared to the $f_{FPGA}/4$ MACs of Solution 1.

1.3 Function Pipeline: Solution 3

The goal of this section of the assignment is to perform further optimizations to implement a design that supports function level pipeline with an II of 1. This means that one complete inner product calculation can be started on consecutive clock cycles.

First create a new solution, *Solution3* in the project workspace. Uncheck the Copy directives and constraints from solution checkbox so that any directives previously defined for a previous solution will not be propagated to this solution. New directives will be specified for this implementation. With the source code `inner_product.cpp` window open, access the Directives window. Right click on the top-level function name *inner-product*, select Insert Directive and following steps similar to those in Section 1.2, add a PIPELINE directive to the function. Run C synthesis and then look at the Performance Profile in the Analysis view. The result should show a function initiation interval of 4. Meaning the function can be invoked every 4 clock cycles. This is a factor of two improvement over the function initiation interval of 8 in Solution2, but not the initiation interval of 1 that is the design objective for this solution. What is going on?

Scroll through the dialog in the console window and you should see the warning.

Try to figure out what caused this warning and timing violation?

Explore achieving function II of 1 by adding memory partitioning directives to each of the a and b and function pipelining.

1.4 Loop Pipelining with Partial Unrolling : Solution 4

So far two useful designs have been produced from a common C code base . The fully pipelined single multiplier sequential inner product circuit in Solution2 and the 8-multiplier fully parallel and task-level pipelined design of Solution3. These two designs represent two extremes in design space. There are intermediate design points that offer alternative tradeoffs between implementation cost and performance as measured in terms of vector operations per second. In this section of the assignment an implementation that is positioned, performance-wise, in between the extremes of the sequential and parallel designs will be generated. This design will use 2-multipliers. Create a new solution, Solution4. The

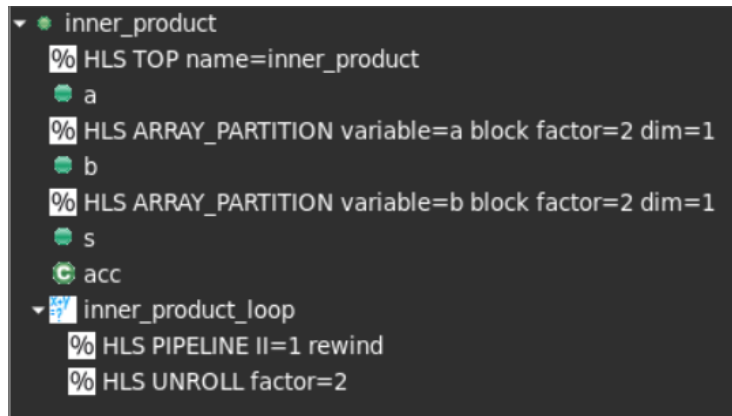


Figure 23: Inner product compute loop pipelined and unrolled by a factor of 2 to generate a design resourced with 2 multipliers and performing 2 multiplications on each clock cycle.

main compute loop *inner_product_loop* is to be pipelined and partially unrolled. Add the directives to the design as shown in Figure 23. This unrolls the loop by a factor of 2 and generates an implementation that is resourced with two multipliers. Since the loop is also pipelined, this means that 2 multiplications are computed on every clock cycle. One inner product is computed every $N/2$ clock cycles. The memories must also be partitioned in order to supply the appropriate number of read ports to supply the required number of operands to the multipliers. In this case each multiplier requires an operand from the a and b memories, so each memory must have 2 ports. Look at the interface description to confirm this is the case.

Run C/RTL Cosimulation to confirm that the RTL generated with this set of directives is functionally correct.

1.5 Loop Pipelining with Partial Unrolling : Solution 5

Make a new solution and create a 4-multiplier design and that executes 4 multiplications on every clock cycle.

Make sure to run C/RTL Cosimulation to confirm that the RTL generated with this set of directives is functionally correct.

2 Assignment Submission

Your assignment submission is to consist of the following components, the submission of the archived project and a small project report which is to include the items listed below.

Use word, or other software of your choice, to generate your report. Don't simply print this page and send me a image of the tables completed by hand,

Table 1: Resource footprint for the 5 implementations of the inner product function.

	solution1	solution2	solution3	solution4	solution5
BRAM					
DSP48					
FF					
LUT					

Table 2: The compute rates of the 5 inner product architectures.

	solution1	solution2	solution3	solution4	solution5
inner_product_loop() pipelined	no	yes	-	yes	yes
inner_product() function pipelined	no	no	yes	no	no
inner_product_loop(); Partially Unrolled	no	no	no	yes, by 2	yes, by 4
MPYs/Sec.					
Vector Ops/Sec.					

take the time to prepare a small report, no hand written work please.

1. Use the Vitis HLS project archive capability to archive and then email your archive to me. A project is archived by selecting, from the HLS IDE ribbon menu, File / Archive Project. Uncheck the Active Solution Only checkbox to ensure that all of the project solutions are added to the archive and check the Include Run Results checkbox so that all of the synthesis and implementation results are included in the archive. Save your archive file using the naming convention first_name.lastname.elen226.a1. The archive .zip file is saved in the project directory.
2. For the 5 solutions you generated in this assignment complete the following table and include it in your project report. Take the time to generate a report and build your own version of the tables below.
3. For the 5 solutions you generated complete the following table and include it in your project report. Use the LUT, FF and DSP48 resource count from the place-and-route runs rather than the estimates generated from C Synthesis. To complete the entries for the final two lines in the table, *MPYs/Sec.* and *Vector Ops/Sec.*, assume that the FPGA clock frequency is 250 MHz.